

Book Proposal

Into the data: learning from machine learning

Adrian Mackenzie Sociology, Lancaster University Bailrigg, LA14YD, UK

Overview

The key question isn't 'How much will be automated?' It's how we'll conceive of whatever *can't* be automated at a given time (Lanier 2013, 77).

This book sets out to construct some ways of critically engaging with data that neither blithely affirm beliefs in the power of data, nor reject beliefs in data as pure hype, nor quails in the face of an incomprehensibly hypercomplex technological revolution. It focuses empirically on certain power-laden practices and techniques at the heart of contemporary media, sciences, government, commerce and industry: *machine learning*. Machine learning is widely used in industry, science, commerce, media and government to program computers to find patterns, associations, and correlations, to classify events and make predictions on a large scale. The Microsoft Kinect motion-sensing device is a mundane example. It uses a machine learning technique called 'decision trees' to identify and classify gestures and body movements, and these gestures, poses and movements become ways of playing a game. Similar predictive and classificatory mechanisms appear in many gadgets (for instance, the face recognition feature in many digital cameras). Almost any search engines' results are shaped by machine learning algorithms, as are many of the recommendations and suggestions generated by social media platforms. Standard scientific applications include detection of abnormal tissue growths in medical scans or the classification of biological function in whole genome sequencing projects. Machine learning, as a set of techniques for classifying and predicting, is widely heralded in the form of data mining, predictive analytics or knowledge discovery as a vital component of contemporary innovation and economic growth. Machine learning is heavily used in search engines, social

network media, high-frequency trading, in increasingly data-intensive scientific practices in astrophysics, genomics, ecology or material science, in manifold devices and control systems, and of course, by government intelligence and security agencies in massive data surveillance programs.

There is intense technical interest and investment in these techniques, and they have developed rapidly in the last two decades. While many commentators acknowledge the importance of algorithms, machine learning techniques have hardly been discussed in humanities and social science literature, even in digital humanities, which heavily relies on them (for instance, in the growing use of 'topic models'). Examining key machine learning techniques and practices drawn from social network media, finance markets, image processing, robotics, and contemporary sciences such as genomics and epidemiology, *Into the data* does not exhaustively describe who does machine learning, where and how. It asks how and why such techniques have moved or circulated so widely, and how their wide circulation might be understood. It suggests that the growth of machine learning does signal changes in modes of knowing and acting, and these changes might imply altered modes of thought for the humanities and social sciences. Seeking to make sense of these changes, *In the Data* is a self-reflexive experiment in writing for humanities and social science audiences that combines code, data and diagram, text, and number with the goal of materially imagining doing machine learning differently. In both analysing and re-telling techniques found at the intersection of contemporary sciences, business, government, network and digital media, the book attempts to both make potent data practices more visible, and to facilitate greater overlaps and entanglements between various science and political, economic and cultural processes associated with data.

Key aims for the book are:

- The modes of existence of data (in the forms of 'big data,' 'open data,' the rise of 'data analytics' and 'data science') should be analysed critically by situating them

in relation to specific settings, techniques and practices. These settings, techniques and practices have complex genealogies, criss-crossing sciences, industries, military, commercial and governmental domains. While critical accounts and indeed skepticism about the value of data are appearing, the provenance and distribution of data practices such as learning algorithms, predictive modelling and classification needs much more critical de-ciphering and empirical attention. This book focuses on the role of machine learning (or data-mining, as it is called in some domains) because the dynamic and varied life of these techniques themselves have received least attention of all. So much depends on and is decided by the models, algorithms and techniques of machine learning, yet they are hardly ever discussed in their own right.

- Humanities and social science responses to data techniques should be methodologically and conceptually inventive, and include appropriation and re-purposing of the techniques and practices. This is major undertaking for the book. It seeks to a broad-ranging way of thinking about data, and what is 'in data' that both soberly appraises beliefs about data, and offers ways of evaluating what is at stake in data as it is processed, explored, organised, filtered and classified by machine learning techniques. The question here is: what can humanities and critical social sciences learn from machine learning?

Approaches used in the book

A broad ethico-political concern underpins *In the Data*. Much contemporary data practice is closely allied to the predictive ambitions of business, the military and states, as well as sciences and media. The recent upswing in data talk continues and intensifies the technoscientific 'Regime of Computation' (Hayles 2005). It is no accident that autonomous military vehicles, large-scale analysis of sentiment social media for commercial or security purposes, or face recognition for national border control are iconic examples of machine

learning in action. A key question for critical humanities and social science researchers, as well as activists, non-government groups and civil society actors of many kinds, is how to make sense of such data practices. They are hard to render visible since they take place largely on platforms that are not publicly accessible. Rendering such practices visible, learning to track their workings, and inventing different ways of working with them: these concerns lie at the core of the analytical and experimental writing practices of *In the Data*.

Broadly speaking, the writing seeks to respond to the long-standing call for what in a widely cited passage Donna Haraway more than a decade ago termed 'diffraction': 'What we need is to make a difference in material-semiotic apparatuses, to diffract the rays of technoscience so that we get more promising interference patterns on the recording films of our lives and bodies' (Haraway 1997, 17). There are a growing number of attempts to adapt and reinvent data practices such as machine learning for less overtly biopolitically laded, security-minded or explicitly commercially-motivated purposes (the growth in university 'data science' courses might be one example (Schutt and O'Neil 2013); see the 'OccupyData' group in New York, N.Y. for one such example (OccupyData 2013); many citizen science projects have something of this flavours to them too). Some of these will be discussed in the course of the book. But the core issue is whether there are forms of thought implicit to machine learning in all its variety that potentially support different ways of thinking about agency, value, power, experience and subjectivity today.

In order to bring data, code, images and text together diffractively, the book draws on some 'executable paper' formats developed in recent scientific publishing. It mingles code written in [R](#), an important statistical modelling, data manipulation and visualization programming language, with code written in Python and Javascript, two of the most popular general programming languages in use today. Some of the empirical materials in the book have been garnered, ordered, analysed and displayed using R, Python and Javascript. More importantly, since all machine learning techniques are implemented in code, working with code is an important way of navigating and exploring the architecture of these

techniques. Code allows movement, visualization and demonstration of machine learning in practice. Code excerpts form part of the text, and will be the object of commentary and analysis, alongside diagrams and graphs generated by the code. All of the code will be available at a public code repository (github.com).

The motivation for the executable format of this book is partly ethnographic and partly experimental. While somewhat agnostic about the ideal of reproducible or executable research, the possibility of combining code, text and image offers scope for different ways of writing and thinking as self-reflexive social researchers. A long line of ethnographers have learned to do what they are observing (as in 'observant participation' (Wacquant 2004)). This has include working in factories, going to prison, spending time in isolated, far-flung or ostensibly boring places, learning techniques ranging from weaving and cooking to playing the piano or programming robots. Ethnographic presence in a particular setting is normally documented through text, photographs, diagrams and occasionally film or audio recordings, and aims to make sense of this setting in ways that both resonate with the people who live there and with people who don't. The forms of observant participation in this book include attending machine learning and other data-related courses (online and in class rooms), competing in online machine learning competitions, reconstructing and implementing algorithms, building predictive models and visualization as a way to present machine learning. These forms of participation are documented by treating coding as both an object of analysis and as part of the ethnographic writing process. Hence forms of data practice analysed in the book become, to varying extents, a mode of writing the book. Several versions of this recursivity will appear in the chapters.

The experimental character of this writing entails both practical and theoretical challenges. Practically, the book works with a range of code constructs, some key mathematical formulae as well as data tables and data graphics. They are not typically found in humanities and qualitative social science writing, although they are extremely common in many scientific fields. The presence of code, formulae and graphics in *In the Data* is not meant to instruct readers in machine learning algorithms or statistical inference. Accompanied by

forms of explication and commentary, they are intended to allow readers to pay close attention to the forms of thought or contemporary equipment (Rabinow 2003) at work in the manifold data practices of sciences or business analytics, and to begin to borrow, appropriate and re-purpose some of the patterns of thought for different purposes. The theoretical ambition here is to treat coding-writing as a way of constructing concepts, metaphors and ways of speaking about contemporary entanglements of subjectivity and computation.

The architecture of the book

The book is organised around two different axes:

1. On one axis, the 'technique axis,' the chapters of the book catalogue, document and analyse some of the most widely used machine learning techniques of working with data (Hastie, Tibshirani, and Friedman 2009). As mentioned above, the diverse techniques analysed on this axis -- linear and logistic regression models, decision trees, clustering algorithms, neural networks, support vector machines and Markov Chain Monte Carlo simulation -- are used across scientific, industrial, biomedical, commercial and military settings. Their extraordinary success in populating these domains cannot be explained in terms of IT or digitisation in general. Using case studies, pedagogical and instructional materials, scientific papers, software libraries and sample datasets, the chapters explore how these techniques, and their implementation as 'learning algorithms,' rely on widely shared assumptions about the problems of knowing, acting, responding to or predicting how things happen. To the extent that a situation can be reshaped to conform to these assumptions, these techniques gain traction. The primary empirical materials come from the vast scientific and engineering research literature generated over the last 50 years or so. Around 100,000 references to papers and books, many of them only slightly different and perhaps little read, bear the traces of widening, overlapping circulation of actions, narratives, and problems associated with techniques and practices of modelling and

classifying data.

2. The other axis of the book is 'recursive reconstruction.' Along this axis, chapters of the book engage with the messiness, complications, and frictions of working with datasets, with predictive models and forms of visualization ranging from standard plots of curves to network graphics. The diagrams, functions and code constructs arrayed along this axis are sometimes drawn from various scientific or commercial settings, but always with an eye on how working with data impinges on some aspect of experience, sociality or feeling, and how these feelings for data might lead to new forms of responsiveness and respons-ability. It develops the proposal that specific situated entanglements of subjectivity and data practice might lie, and to open up different ways of thinking about contemporary experience as it is increasingly pervaded and subtly (or not subtly) modulated by data-driven processes. The reconstruction of data practices draws on the pragmatist philosopher John Dewey's notion of philosophy as an empirical reconstruction of experience (Dewey 1957; Dewey 2004). The kinds of experience reconstructed range from encounters with senses of agency, proximity, and otherness, but also with practices of work, sensations of familiarity, strangeness, as well as feelings such as recognition, inclusion or indifference. These reconstructive moves will be linked to broader debates around politics, ethics, publics, democracy, power, equality and differences.

Existing academic literature and framing of the book

The existing literature relevant to this monograph come from a variety of disciplines. The critical work on data is largely found in science and technology studies (STS), and some parts of information science. Software studies and anthropological accounts of software cultures are highly relevant in reading machine learning algorithms and data visualization software. A broader theoretical background here includes recent reappraisals of pragmatism (particularly William James, C.S Peirce), feminist and other work on materialities,

as well as strands of largely European contemporary philosophy relating to experience, space-time, science, calculation and events. A final reference point comes from recent attempts in social sciences and humanities to reinvent methods of research.

In STS, work on calculation (Callon and Law 2005), data practice (Edwards et al. 2011), databases (Bowker 2005) and digital data more generally (Latour et al. 2012) have extensively discussed how science assembles numbers, observations, instruments, readings and databases. This work forms an important part of the background of this book since machine learning has heavy mathematical underpinnings and institutional dynamics. The STS work has broadly re-theorised many different aspects of data, ranging across collection, measurement, calculation, archiving, labelling and visualising. Much of this work is based on ethnographic case studies of laboratories, technical devices, standards and controversies. It has notably developed ways of analysing its objects relationally (as in actor-network approaches), and with an eye on entanglements and hybridisation of human and non-human entities. While *In the Data* is not by any means a standard laboratory ethnography, it does rely on practices of participant observation and analytical approaches found in STS.

The history of statistics, number and mathematics also frame important aspects of *In the Data*. Works such as Theodore Porter's *Trust in Numbers* (Porter 1996), Lorraine Daston's work on probability (Daston 1988), or Alain Desrosiere's *The Politics of Large Numbers* (Desrosieres 1998) amongst others not only provide background for many of the statistical techniques used in machine learning, they suggest that numerical data and numbers have had an eventful course of development from the 18th to the 20th century. While much of this historical work leaves just around the time when machine learning approaches are emerging (1960-1970s), it provides an extremely useful way to contextualise key traits in the contemporary data practice, ranging from genres of visualization to underlying concepts of probability, chance or error.

The nascent field of software studies has begun to develop ways of analyzing software and

code, ranging from source code files to large assemblages. Coupled with media studies and media archaeology-type approaches, software studies has developed genealogies, critical framings and methods of reading many different aspects of software. Work in this field ranges from quite high-level analyses such as Wendy Chun's *Programmed Visions* (Chun 2011) or Lev Manovich's *Cultural Software* (Manovich 2011), or Alex Galloway and Eugene Thacker's work (Galloway and Thacker 2007) through to studies of specific code objects (as for instance in many of the entries in the *Software Studies: A Lexicon* volume (Fuller 2007)) or analysis of code as speech (Cox and MacLean 2012). Other work should be included here (along with related work on 'platform studies'), but for present purposes, the key influence of software studies consists in its treatment of software, computer code, algorithms and protocols as first-ranking objects of social and cultural analysis. Some literature on data flow and data visualization has started to appear, but machine learning doesn't figure in it (Beer and Burrows 2013). A broader background of work on software cultures (C. M. Kelty 2008; Coleman 2009), with their important re-thinking of publics, property and value, is also relevant, but differs greatly in its emphasis on software production and social movements.

A broader range of theoretical approaches informs this book. These include on events, materiality, experience, and capitalism from scholars that include Brian Massumi on radical empiricism (Massumi 2000), Nigel Thrift on time-space signatures of calculation (Thrift 2005), Celia Lury on topological conceptions of culture (Lury, Parisi, and Terranova 2012), Manuel Delanda on simulation in philosophy and social science (DeLanda 2002; DeLanda 2011), Anna Munster on conjunctive experience in networks (Munster 2013), or Luciana Parisi on the contagiousness of computation (Parisi 2013). Many of these authors share an interest in re-thinking notions of experience, body, event, time-space and materiality in the context of ongoing transformations of media and technology. Many of them draw on philosophers such as William James or A.N. Whitehead to question taken-for-granted concepts of nature, life or agency. Again, this loose coalescence of work cannot be adequately summarised or even limned here, but it indicates something of

the theoretical registers on which *In the Data* will work.

The final framing body of work is even less coherent, but nevertheless important: it largely comprises threads of research and debate about methods today in social sciences and humanities. This literature tends to treat the growth of digital data as both posing a problem and an opportunity for research in social sciences and humanities. The problem, as framed by sociologists such as Andrew Abbott (Abbott 2001) or Mike Savage (Savage 2009), is that existing quantitative methods in social science cannot match the efficacy of quantitative methods in the natural or applied sciences, nor those used in business and marketing (e.g. as in analysis of transaction data). Some social scientists advocate the development of 'computational sociology' (King 2011). A version of the same crisis can be found in digital humanities, and has prompted developments such as 'cultural analytics' (Manovich 2009). Commonly, these responses advocate a pattern-based approach to working with social or cultural data, and in this respect, they mirror some of the commitments in machine learning to finding the function that generates the data. Whilst very sympathetic to and in some ways aligned with these debates, *In the Data* also aims to offer something other than a set of 'better' data methods. It is more closely aligned with work that seeks to re-think social science methods in the light of new flows of data (Marres 2012) and its temporalities (Uprichard 2012).

Readership and market

The readership for the book is quite diverse, since data practices and indeed machine learning itself are of interest to a growing audiences. One set of readers for the book come from disciplines such as sociology, anthropology, media and cultural studies, and social geography who are grappling with the promise of data both as an object of analysis and in terms of a transformation of their own ways of researching. Another set of readers for the book come from the burgeoning 'data science' courses being offered in North American, UK, SE-Asian/Pacific, and European universities. While these courses are largely

focused on techniques of organising, visualising and modelling data, many of them are also open to thinking about the transformations in knowledge and value associated with contemporary data practice. The book is written very much with these kind of readers in mind. It will minimize reference to social theory in order to maintain more accessibility to these readers. While I am keen to keep the social and cultural theory side of the book in the margins, the book will introduce some technical terminology, and indeed some mathematical formulations. But this technical material will be analysed rather than assumed as background.

Timetable

All of the chapters exist in draft form, or as conference papers. Writing an introduction, conclusion, and revising the drafts will take roughly 12 months.

- draft conclusion: 1 month
- draft introduction: 1 month
- draft chapter 4: 2 months
- revise chapter 3,4,5,6,7,8 drafts: 3 months
- revise chapter 4: 1 month
- revise whole manuscript: 3 months

Format of the book

The book has a standard chapter format. It will include several dozen code-generated figures, diagrams or plots, as well as a number of tables. The Python and R code, and datasets used to generate these components of the text will be available through the public code repository github.com. The Markdown text of the book will be also part of this code repository. Electronic versions of the book will display colour versions of the plots, and be hyperlinked to both the code-data components on github, and to various relevant

URLs. The predicted wordcount is 85,000 - 90,000 words. It will include approximately 20 diagrams or graphics.

Chapter outline

introduction: Into the Data

The introduction will begin with several relatively familiar examples of machine learning drawn from a variety of fields over the last decade or so -- handwriting recognition, face recognition, autonomous robots (Thrun et al. 2006), credit card checks, and cancer prognosis. It will highlight these examples as symptoms of the wide-ranging and often long-standing investments in knowledge, control, prediction and decision-making associated with data flows. At the same time, it will suggest how these tracking some of the transformations might elicit changes in how humanities and social science researchers understand their own work.

These examples will also provide a preliminary overview of the techniques of machine learning discussed in the book -- supervised and unsupervised learning, the differences between classification, regression, and clustering and important notions such as learning and prediction. They will also highlight contrasts between disciplines such as computer science and statistics that develop machine learning techniques, as well as illustrate the overlaps between data-mining, pattern recognition, knowledge discovery, artificial intelligence, information retrieval, signal processing, machine learning etc. Practically, these examples will also implicitly present some of the methods used in the subsequent chapters, including the role of databases, data structures, code constructs, diagrams, and algorithms in typical scientific and industry practices of modelling.

These examples will also stage some of wider questions in the book about the promise of data. These include the oft-mentioned 'end of theory' prediction (Chris Anderson, *Wired* magazine, 2008), and the many claims and controversies about data analytics, machine learning and the 'power of big data' in physical, life and social sciences, in business, gov-

ernment and industry. Claims about power of data, and responses to these claims -- ranging from downright skepticism to enthusiastic embrace -- will be discussed here with an eye on what these debates about data mean for research practices in the social sciences and humanities themselves in terms of their topics of research and how they do research.

Finally, the introduction will sketch the themes of 'in the data' and 'modes of machine thought,' drawing on a range of work drawn from pragmatist philosophers such as C.S. Peirce (abduction and diagrams), William James on experience (James 1996), John Dewey on 'reconstruction' (Dewey 1957), Alfred N. Whitehead on 'abstraction' (Whitehead 1958) and from recent social and cultural theory such as Isabelle Stengers on experiment (Stengers 2008); Gilles Deleuze & Felix Guattari on scientific functions, and (Deleuze and Guattari 1994); Celia Lury on topology (Lury, Parisi, and Terranova 2012)). In order to contextualise forms of data thought, the introduction will also sketch some points of departure drawn from software studies work on algorithms and databases, science studies work on calculation, statistics, number, device, image and diagram, as well as accounts of subjectivity, experience (Berlant 2007) or (Adams, Murphy, and Clarke 2009) and materiality cross-cutting all of the above. This spectrum of work from across disciplines provide scaffolding and departure points for much of the book.

Part I: Data, Functions and Forms

The four chapters of Part I explore underlying the major underlying spatial, ordering and counting practices of the many different techniques comprising machine-learning, pattern recognition and data-mining. It seeks to show how these practices diverge and multiply across a range of settings, and how they coalesce around a set of generic intuitions of difference that are both powerfully generative yet highly constrained.

1. Writing about data

Key techniques: linear models, perceptron This chapter is primarily a methodological discussion that addresses several different problems in working with machine learning. These problems range from quite philosophical issues through to quite practical ones. At the philosophical end, it poses some basic problems in responding to science and technology (for instance, the anthropologist Paul Rabinow's work on the concept of 'equipment' (Rabinow 2003); the philosopher A.N. Whitehead on the spatial dimensionality of thought (Whitehead 1958); the philosopher Anne-Marie Mol's notion of praxiography -- writing about practices (Mol 2003)) to discuss how we might think about working on highly technical or scientific areas such as machine learning in ways that allow consideration of their more general situation. It also draws on some cultural and psychoanalytic accounts of architecture and objects (Bollas 2009; Wilson 2010) to suggest how the researcher him or herself relate to objects of research. Finally in this vein, the chapter poses the methodological problem of working with large bodies of scientific and technical literature, and illustrates this by describing the growth of machine learning techniques in science and engineering research since the early 1960s, focusing on the growth of key techniques and algorithms (C. Kelty and Landecker 2009). Oriented by these questions about thinking, techniques, subjectivity and literature, the chapter then presents a series of vignettes that display some of the ways in which research and writing critical accounts of data cultures and data economies can make use of the tools, techniques, instruments and services of 'data science' to generate textual, diagrammatic and modelised accounts of contemporary culture. A standard teaching example -- house-price prediction -- links these vignettes, but the real focus here is on two foundational issues: how machine learning treats data as a dimensional material that it seeks to reshape or recast in different dimensions ('models'); how implementing machine learning techniques shifts our relation to them. The chapter describes some of the transformations in software, network and scientific cultures that underpin the recent growth in data techniques and methods. These range across transformations in statistical science associated with greater computational capacity; the mutations

in network, database and digital device architectures and infrastructures that yield much greater abundance of data in various forms; and the intermeshing of knowledge economies with the media, communication, transaction, transport and logistics systems. It will trace how the lateral associations and multivalencies of data have developed through key software artefacts such as the widely used R programming language, and in generic programming languages such as Python. Reflecting on the author's own history of working with machine learning or online accounts of machine learning, well as the ecology of thousands of software packages associated with the statistical programming language R, the aim here explore some of the material transformations that might open machine learning to wider engagement.

2. Machines finding functions

Key techniques: logistic function, cost function, gradient descent The *learning* in machine learning is often conceived as a form of function-finding. Developing ways to read functions, and to highlight the ways they configure differences, similarities, proximities, distances, boundaries and changes, is key to both understanding how machine learning works, and also to imagining ways in which the abundance of functions can be understood. This chapter examines the proliferation of predictive models, classifiers, clustering and other analytic devices in terms of *functions*. Drawing on statistical machine learning texts (Hastie, Tibshirani, and Friedman 2009), and more philosophical accounts of functions (e.g. (Deleuze and Guattari 1994; Whitehead 1958)), the chapter introduces several key instances of the function in machine learning, shows how functions underpin the generation of curves, and how movement along lines, curves and across planes. While later chapters will range across a wider variety of mathematical functions and forms, this chapter will focus on basic components of many machine learning techniques: a function that maps a state of affairs onto a general abstract form (usually somewhat Euclidean - lines, planes, regular curves), a function that measures errors or disparities, and a function that optimises the mapping between a specific abstract form and the state of affairs em-

bodied in data. It will exemplify these components through simple examples such as the logistic function, the least-squared error cost function, and the gradient descent algorithm. Finding functions, the chapter will suggest, is the central way in which machine learning connects experience to abstraction. Standing back from mathematical-algorithmic techniques, this chapter suggests that finding the functions that inscribe lines and surfaces in data is a powerful form of imitation that tends to remake the world in certain ways. This re-making may be inimical to social life, but it is possible to observe it, and finding ways to observe it might foster different responses to it.

3. Finding many patterns in data

Key techniques: decision trees, support vector machines & random forests This chapter compares three of the most popular machine learning classification techniques: decision trees, random forests and support vector machines. For the last decade, one of the best-performing 'off-the-shelf' machine learning algorithm has been a technique known broadly as 'support vector machines' (SVM; see (Vapnik 1999)). Versions of the decision tree have since the early 1970s been heavily used across science, business, media and medicine to classify and predict differences. Finally, random forests massively multiply decision trees to overcome problems associated with much greater scale and diversity of data. The chapter examines the spatial operations of these widely used algorithms both against the background of a spectrum of other statistical machine learning techniques, and more importantly, in terms of the *forms of movement* they bring to data practice. The key issue here is how machine learning moves through data, and how it copes with the expanding dimensionality of data. These models are less Euclidean or geometrical in their movement. Since the 1950s, scientists have been aware of the 'curse of dimensionality' (Bellman 1961), which arises when the dimensions of the data increase. Algorithms such as decision trees, SVM, and random forests introduce different ways of navigating data dimensionality. While lines, curves and planes, as discussed the previous chapter, inscribe boundaries and geometrical order in data, the advent of increasingly extended and partic-

ularly 'wide' datasets (many variables) engenders models that embrace high-dimensional spaces in which patterns are much more unstable and somewhat virtual.

4. Believing in the strength of numbers

Key techniques: Naive Bayes, Markov Chain Monte Carlo simulations, Bayesian networks; The topic in this chapter is the role of randomness, chance and number in machine learning. Machine learning techniques are suffused with probabilistic modes of thought. This chapter foregrounds probability and the changes in probability associated with both statistical pattern recognition models such as the Naive Bayes classifier (often used to filter spam email) and the so-called 'Bayesian revolution' in statistical practice that took shape in the early 1990s in particular in the form of Markov Chain Monte Carlo simulation (MCMC). The computationally intensive techniques of Bayesian analysis treat all numbers as potentially random variables; that is, as best described by probability distributions whose interactions need to be carefully explored. By contrast, the probabilistic machine learning models exemplified by Naive Bayes treat numbers as if they have little relation to each other. The chapter traces two important implications of this contrast between ways of working with probability. The popularity of MCMC is a striking example of a technique moving across fields, and the chapter will trace some of the ramifications of the heavily-used MCMC technique in fields ranging from nuclear physics, image processing to political science and online gaming. Second, although they both treat potentially all important numbers as a matter of probability calculus, the contrast between computationally intensive, MCMC and probabilistic models such as Naive Bayes suggest very different beliefs in the power of computation. A broader question here will be framed by reference to notions of expectation and belief: what mode of belief in probability better sensitises us to what machine learning or pattern recognition models do in given situations?

Part II: Problems with Data

The three chapters of Part II explore what happens as the major intuitions of machine-learning encounter things, events and people. The chapter deal with transformation in scale, diversity and work.

5. What does machine learning do to data?

Key techniques: decision trees, random forests, self-organising maps Problems of scale have framed many aspects of machine learning. The increasingly large scale of data motivates automated forms of analysis. The problems of discerning small differences motivates attempts to gather data on a larger scale. Genomics is a provocative form of data thought in several respects. It relentlessly treats one type of quite flat or mono-dimensional data -- nucleic acid sequences or 'base sequences'-- as the key to potentially subtle or complicated biological processes in all their plasticity and mutability. While it is not at all clear that this treatment will be effective, it has animated efforts to find shape or pattern in sequence data that stand as a limit case for data-driven research more generally. At the same time, genomics is a scientific discipline almost overwhelmed by the effectiveness of its own instruments in generating data. Sequencing in contemporary genomics (that is, post-Human Genome Project and after the advent of so-called 'high-throughput' or 'next generation sequencers'; this is roughly 2007 onwards) is perhaps the most concerted effort to count things every undertaken in the life sciences. The rate of production of sequence data from next generation sequencers exceeds Moore's Law, the standard 18-24 month doubling time for the number of transistors in an integrated circuits. This sequence data needs to be stored and analysed in rhythms that differ from many other settings where the growth of data can be managed through more memory and computer processing speed. Third, genomic researchers have been extraordinarily adaptive in positioning their work on the borders of cutting edge infrastructure development, machine learning and data-mining, and the life sciences. Genomics (and bioinformatics) loom large in the machine

learning literature itself since the mid-1990s. The flatness of sequence data has been heavily leveraged by this positioning. The biological objects of genomics - genomes -- have been progressively transformed and re-shaped in ways that might be instructive for data more generally. This chapter explores how machine learning has imbued genomes with an increasingly topological character (and particularly, the growth of 'topological data analysis' (Carlsson 2009; Singh, Mémoli, and Carlsson 2007) as well as the topological turn in culture (Lury, Parisi, and Terranova 2012)), and practically, with the rich ecology of programmatically accessible bioinformatics tools and archives that on the one hand permits sequence data to move relatively freely (especially in comparison to much commercial or even social media data), but on the one hand poses question as to who wants or needs the data.

6. Are there enough numbers in world?

Key techniques: transmission models, nested models A predominant narrative around data in many contemporary settings urges that more data makes all problems solveable. This narrative is usually accompanied by an 'abundance of data' ('big data', 'data deluge', etc) narrative, in which the advent of data corresponds to a groundswell change in how we make sense of and intervene in events. Versions of these narratives surface in genomics, business analytics, and infrastructure management (e.g. in smart energy grids), as well as crisis-events such as financial collapses or epidemics. Via a case study of different data flows during the 2009 A/H1N1 'swine flu' epidemic, this chapter develops an alternative narrative of data flow in terms of number supply chain logistics. The chapter reconstructs a real-time epidemiological model that combines clinical reports, laboratory test data, web surveys, urban population mixing patterns in order to disentangle biological and social forms of contagion and infection during the 2009 epidemic in London. In reconstructing this model, a model that is typical in complicated engagement with numbers of diverse origins, the chapter will suggest that the largely homogeneous data flows envisaged and embraced in many forms of data practice largely ignore the problem of the interactions

between different agents. It specifically contrasts the much publicised Google Flu Trends approach to 'flu prediction, which is based on search query volumes, with epidemiological models based on multiple forms of surveillance data. The chapter argues that data practices during crises or times of great uncertainty, entail hybrid integrations of existing data practice and new forms of data. This reconstruction challenges us to re-evaluate how we think about numbers. By following some of the ways numbers circulate through the epidemiological models, we can discern to a semiotic-material faultline running through contemporary number formations. Numbers semiotically and materially embrace both events and degrees of belief. If numbers are crucial in the data economy, then instabilities in their mode of existence will affect much of what happens to data. While much of the machine learning taking place in commercial and operational settings is decidedly non-Bayesian, the popularity of MCMC and Bayesian approaches in contemporary sciences suggests a tension in what counts as number.

7. Optimising machine learners

Key techniques: ensembles, RandomForest The chapter focuses on the forms of work and experience associated with contemporary data practice, situated within plural data and knowledge economies. Software developers, hackers, statisticians, 'data scientists,' as well as social scientists, are changed by forms of data thought. The case study in this chapter is data prediction contests run by the [Kaggle.com](https://www.kaggle.com) as well as academic-based competitions. In these competitions, competitors from diverse technical and geographic backgrounds compete to construct predictive models for specific datasets -- the Netflix recommendation competition; the Facebook 'find a friend' competition; or the Titanic survivor problem -- using whatever machine learning techniques they can bring to bear. These competitions, conducted on web-based platforms, are useful ways to track contemporary data practices. Combined with some examples of presentations by academic researchers (for instance, Stanford University's Andrew Ng whose YouTube lectures have attracted 100,000s of views), industry conferences (for instance, at the annual Predictive Analyt-

ics World events), this chapter will track the kinds of technical and affective investment associated with popular data modelling techniques such as Random Forest. It is possible, I will suggest, to read a technique as a partial subjectification, in that it affects how they experience and materially engage with data. In order to apprehend the character and texture of these subjectifications, the chapter links university research, commercial and non-commercial adoption, and flows of technical expertise. Again, this chapter has some auto-ethnographic vignettes, as the author has participated in some of these competitions.

Conclusion: Out of the Data

The conclusion draws together the main threads running through the previous chapter, and sets out a series of questions and provocations for thinking with data. The conclusion will stand back from fine-grained empirical approach to data and data practice adopted in the preceding chapters in order to think more about we -- social scientists, humanities scholars -- might invent or create in the midst of data. While this book has a critical angle to it (so many claims about and beliefs in data plainly deserve critique for their conservative and naive approach to things), it is principally concerned with conceptual invention through doing things with data. The work of learning about machine learning, and learning about it in a way that is deeply embodied or practically embodied, brings with it altered ways of thinking about, questioning and integrating what is happening to data more generally. It highlights the key argument that has run through the book about the plural dimensionality of data as it is aggregated, tabulated, summarised and modelled in contemporary data and signal processes, and as well as the extraordinary mobility or kinetic energy of generic machine learning methods. In discussing the shifting dimensionality of data, and the kinetics of methods, the conclusion will attempt to sketch out how some promising ways of thinking with data might proceed.

References

Abbott, Andrew. 2001. *Time matters: on theory and method*. University of Chicago press.

Adams, Vincanne, Michelle Murphy, and Adele E. Clarke. 2009. "Anticipation: Technoscience, life, affect, temporality." *Subjectivity* 28 (1) (sep): 246--265. doi:10.1057/sub.2009.18. <http://www.palgrave-journals.com/doi/10.1057/sub.2009.18>.

Beer, David, and Roger Burrows. 2013. "Popular culture, digital archives and the new social life of data." *Theory, Culture & Society*. <http://tcs.sagepub.com/content/early/2013/04/15/0263276413476542.abstract>.

Bellman, Richard. 1961. *Adaptive control processes: a guided tour*. Vol. 4. Princeton university press Princeton. <http://www.getcited.org/pub/101191710>.

Berlant, L. 2007. "Nearly utopian, nearly normal: Post-Fordist affect in La Promesse and Rosetta." *Public Culture* 19 (2): 273. <http://lucian.uchicago.edu/blogs/politicalfeeling/files/2007/10/berlant-la-promesse.pdf>.

Bollas, Christopher. 2009. *The evocative object world*. Taylor & Francis. http://books.google.co.uk/books?hl=en&lr=&id=Hx5T_TUzyh8C&oi=fnd&pg=PP1&dq=christopher+bollas+evocative+object&ots=U30i_rj3I7&sig=N1CW3bEizpUaxTH41IAfDfAbQi4.

Bowker, G. C. 2005. *Memory practices in the sciences*. MIT Press Cambridge, MA.

Callon, M., and J. Law. 2005. "On qualculation, agency, and otherness." *Environment and Planning D* 23 (5): 717.

Carlsson, Gunnar. 2009. "Topology and data." *Bulletin of the American Mathematical Society* 46 (2): 255. <http://www.ams.org/images/carlsson-notes.pdf>.

Chun, Wendy. 2011. *Programmed visions: Software and memory*. The MIT Press.

Coleman, Gabrielle. 2009. "Code is speech: Legal tinkering, expertise, and protest among free and open source software developers." *Cultural Anthropology* 24 (3): 420--

454. <http://onlinelibrary.wiley.com/doi/10.1111/j.1548-1360.2009.01036.x/full>.

Cox, Geoff, and Alex MacLean. 2012. *Speaking Code: Coding as Aesthetic and Political Expression*. Cambridge MA: MIT Press.

Daston, Lorraine. 1988. *Classical Probability in the Enlightenment*. New Brunswick, N.J.: Princeton University Press.

DeLanda, Manuel. 2002. *Intensive Science and Virtual Philosophy*. London & New York: Continuum.

-----, 2011. *Philosophy and simulation: the emergence of synthetic reason*. Continuum. http://books.google.co.uk/books?hl=en&lr=&id=F5wvXkJwFwkC&oi=fnd&pg=PP7&dq=manuel+delanda&ots=_ovoxIUGer&sig=oHqP0fzg6aB6Ydq3EBGYJul_U18.

Deleuze, Gilles, and Félix Guattari. 1994. *What is philosophy?. European perspectives*. New York ; Chichester: Columbia University Press.

Desrosieres, Alain. 1998. *The Politics of Large Numbers: A History of Statistical Reasoning*. Cambridge, Mass: Harvard University Press.

Dewey, John. 1957. *Reconstruction in Philosophy*. 1st ed.. Boston: Beacon Press.

-----, 2004. *Essays in experimental logic*. Mineola, N.Y.: Dover Publications. <http://www.loc.gov/catdir/enhancements/fy0618/2004052766-d.html>.

Edwards, P. N., M. S. Mayernik, A. L. Batcheller, G. C. Bowker, and C. L. Borgman. 2011. "Science friction: Data, metadata, and collaboration." *Social studies of science* 41 (5): 667--690. <http://sss.sagepub.com/content/41/5/667.short>.

Fuller, Mathew, ed. 2007. *Software Studies: a Lexicon*. Cambridge, MA: MIT Press.

Galloway, Alexander R., and Eugene Thacker. 2007. *The exploit : a theory of networks*. Minneapolis: University of Minnesota Press.

Haraway, Donna J. 1997. *Modest_Witness@Second_Millennium. FemaleMan©_Meets OncoMouse™*. New York: Routledge.

Hastie, Trevor, Robert Tibshirani, and Jerome H. Friedman. 2009. *The elements of statistical learning: data mining, inference, and prediction*. New York: Springer.

Hayles, N. Katherine. 2005. *My mother was a computer : digital subjects and literary texts*. Chicago: University of Chicago Press.

James, William. 1996. *Essays in radical empiricism*. Lincoln: University of Nebraska Press.

Kelty, Christopher M. 2008. *Two bits: the cultural significance of free software*. Durham: Duke University Press.

Kelty, Christopher, and Hannah Landecker. 2009. "Ten thousand journal articles later: ethnography of 'The literature' in science." *Empiria: Revista de metodología de ciencias sociales* (18): 173--192. <http://dialnet.unirioja.es/servlet/articulo?codigo=3130617>.

King, Gary. 2011. "Ensuring the Data-Rich Future of the Social Sciences." *Science* 331 (6018) (feb): 719--721. doi:10.1126/science.1197872. <http://www.sciencemag.org/content/331/6018/719.abstract>.

Lanier, Jaron. 2013. *Who owns the future?*. London: Allen Lane.

Latour, Bruno, Pablo Jensen, Tomasso Venturini, S. Grauwin, and D. Boullier. 2012. "The Whole is Always Smaller than its Parts. How Digital Navigation May Modify Social Theory." *British Journal of Sociology* 63 (4): 590--615.

Lury, Celia, Luciana Parisi, and Tiziana Terranova. 2012. "Introduction: The Becoming Topological of Culture." *Theory, Culture & Society* 29 (4-5): 3--35. doi:10.1177/0263276412454552. <http://tcs.sagepub.com/content/29/4-5/3>.

Manovich, Lev. 2009. "Cultural analytics: Visualizing cultural patterns in the era of more media." *Domus* (923). http://softwarestudies.com/cultural_analytics/Manovich_DOMUS.doc.

-----, 2011. *Cultural Software*. \textbackslashCervenec. <http://www.manovich>.

net/DOCS/Manoich.Cultural_Software.2011.pdf.

Marres, Noortje. 2012. "The redistribution of methods: on intervention in digital social research, broadly conceived." *The Sociological Review* 60 (S1): 139--165. <http://onlinelibrary.wiley.com/doi/10.1111/j.1467-954X.2012.02121.x/full>.

Massumi, Brian. 2000. "Too-blue: colour-patch for an expanded empiricism." *Cultural Studies* 14 (2): 177--226.

Mol, Annemarie. 2003. *The Body Multiple: Ontology in Medical Practice*. Durham, N.C: Duke University Press.

Munster, Anna. 2013. *An Aesthesia of Networks: Conjunctive Experience in Art and Technology*. MIT Press. <http://mitpress.mit.edu/books/aesthesia-networks/>.

OccupyData. 2013. "Occupy Data." <http://occupy-data.org/>.

Parisi, Luciana. 2013. *Contagious Architecture: Computation, Aesthetics and Space*. MIT Press.

Porter, T. M. 1996. *Trust in numbers: The pursuit of objectivity in science and public life*. Princeton Univ Pr.

Rabinow, Paul. 2003. *Anthropos Today. Reflections on Modern Equipment*. Princeton and Oxford: Princeton University Press.

Savage, Mike. 2009. "Contemporary Sociology and the Challenge of Descriptive Assemblage." *European Journal of Social Theory* 12 (1) (feb): 155--174. doi:10.1177/1368431008099650. <http://est.sagepub.com/cgi/content/abstract/12/1/155>.

Schutt, Rachel, and Cathy O'Neil. 2013. *Doing data science*. Sebastopol, Calif.: O'Reilly & Associates Inc.

Singh, Gurjeet, Facundo Mémoli, and Gunnar Carlsson. 2007. "Topological methods for the analysis of high dimensional data sets and 3d object recognition." In *Eurographics Symposium on Point-Based Graphics*. Vol. 22. <http://comptop.stanford.edu/preprints/mapperPBG.pdf>.

Stengers, Isabelle. 2008. "Experimenting with Refrains: Subjectivity and the Challenge of Escaping Modern Dualism." *Subjectivity* 22 (1) (may): 38--59. doi:10.1057/sub.2008.6. <http://www.palgrave-journals.com/doi/finder/10.1057/sub.2008.6>.

Thrift, N. J. 2005. *Knowing capitalism*. London: SAGE Publications.

Thrun, Sebastian, Mike Montemerlo, Hendrik Dahlkamp, David Stavens, Andrei Aron, James Diebel, Philip Fong, John Gale, Morgan Halpenny, and Gabriel Hoffmann. 2006. "Stanley: The robot that won the DARPA Grand Challenge." *Journal of field Robotics* 23 (9): 661--692. <http://onlinelibrary.wiley.com/doi/10.1002/rob.20147/abstract>.

Uprichard, Emma. 2012. "Being stuck in (live) time: the sticky sociological imagination." *Sociological Review* 60 (jun): 124--138. doi:10.1111/j.1467-954X.2012.002120.x.

Vapnik, Vladimir. 1999. *The Nature of Statistical Learning Theory*. 22000 ed.. Springer.

Wacquant, Loic. 2004. *Body and Soul. Notebooks of an apprentice boxer*. Oxford: Oxford University Press.

Whitehead, Alfred North. 1958. *Modes of thought; six lectures delivered in Wellesley College, Massachusetts, and two lectures in the University of Chicago*. New York,: Capricorn Books.

Wilson, Elizabeth A. 2010. *Affect and artificial intelligence*. University of Washington Press.