

technique demos

gradient descent

Lesson 2 of Ng; implemented in R [here](#)

Also, useful for demonstration of matrix algebra.

Theory

The model we will get at the end is a line that fits the data, is defined like so:

Setting $x_0 = 1$:

$$h_{\theta}(x) = \theta_0 x_0 + \theta_1 x_1 + \theta_2 x_2 + \dots$$

That can be summarized by (last is matrix notation):

$$h_{\theta}(x) = \sum_{i=0}^n \theta_i x_i = \theta^T x$$

Matrix representation is useful because has good support in software tools.

Goal is to get the line closest to observed data points as possible, thus we can define a cost function that returns the difference of the real data vs myModel:

$$J(\theta) = \frac{1}{2m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)})^2$$

where i is each data example we have and m is their total.

With J we now have a metric to check if the hypotheses line is getting closer to data points or not.

Next step is to find the smaller cost as possible from J , and in fact that's exactly what the [gradient descent algorithm](#) does: starting with an initial guess it iterates to smaller and smaller values of a given function by following the [direction of the derivative](#):

$$x_i := x_{i-1} - \epsilon f'(x_{i-1})$$

Applying to our J :

$$\theta_j := \theta_j - \alpha \frac{\delta}{\delta \theta_j} J(\theta)$$

And doing a bit of calculus on derivatives we get:

$$\theta_j := \theta_j - \alpha \frac{1}{m} \sum_{i=1}^m (h_{\theta}(x^{(i)}) - y^{(i)}) x^{(i)}$$

Where α defines the size of steps of the convergence to θ . Now lets check if all this math really works.




Figure 1: plot of chunk gradient_descent

Implementation - take 1

α is the size of the step taken each iteration

```

alpha = 0.07
m = length(mydata$x)
theta = c(0, 0)
x = mydata$x
y = mydata$y
delta = function(x, y, th, m) {
  sum = 0
  for (i in 1:m) {
    sum = sum + (((t(th) %>% c(1, x[i])) - y[i]) * c(1, x[i]))
  }
  return(sum)
}

# 1 iteration
theta - alpha * 1/m * delta(x, y, theta, m)

[1] 0.07453 0.38002

```

Implementation - take 2

After having a peek at the [Matlab solution](#), I learned that is possible to replace the sum in the equation with a transpose matrix multiplication (like done with the line equation):

$$\theta := \theta - \alpha \frac{1}{m} x^T (x\theta^T - y)$$

So we can get a full matrix implementation:

```

alpha = 0.07
m = length(mydata$x)
theta = matrix(c(0, 0), nrow = 1)
x = matrix(c(rep(1, m), mydata$x), ncol = 2)
y = matrix(mydata$y, ncol = 1)
delta = function(x, y, th) {
  delta = (t(x) %>% ((x %>% t(th)) - y))
  return(t(delta))
}

# 1 iteration
theta - alpha * 1/m * delta(x, y, theta)

      [,1] [,2]
[1,] 0.07453 0.38

```

```

for (i in 1:1500) {
  theta = theta - alpha * 1/m * delta(x, y, theta)
}
theta

      [,1]      [,2]
[1,] 0.7502 0.06388

ex2plot + geom_abline(intercept = theta[1], slope = theta[2])

```




Figure 2: plot of chunk gradient_descent_matrix